

# Elevating R to Supercomputers

Drew Schmidt  
University of Tennessee, Knoxville

July 12, 2013



# Affiliations and Support

The pbdR Core Team

<http://r-pbd.org>

Wei-Chen Chen<sup>1</sup>, George Ostrouchov<sup>1,2</sup>, Pragneshkumar Patel<sup>2</sup>, Drew Schmidt<sup>1</sup>

Ostrouchov, Patel, and Schmidt were supported in part by the project “NICS Remote Data Analysis and Visualization Center” funded by the Office of Cyberinfrastructure of the U.S. National Science Foundation under Award No. ARRA-NSF-OCI-0906324 for NICS-RDAV center.

Chen and Ostrouchov were supported in part by the project “Visual Data Exploration and Analysis of Ultra-large Climate Data” funded by U.S. DOE Office of Science under Contract No. DE-AC05-00OR22725.

---

<sup>1</sup>Computer Science and Mathematics Division, Oak Ridge National Laboratory, Oak Ridge, TN

<sup>2</sup>Remote Data Analysis and Visualization Center, University of Tennessee, Knoxville, TN

# About This Presentation

## Conventions

We use:

- “.” as a decimal mark
- “,” as order of magnitude separator

Example	Yes	No
One million	1,000,000	1.000.000
One half	0.5	0,5
One thousand and one half	1,000.5	1.000,5

# Contents

1 Introduction

2 Benchmarks

3 Challenges

## Why R?

## Why R?

- 1 Because.

## Why R?

- 1 Because.
- 2 R community has growing data size problem.

## Why R?

- 1 Because.
- 2 R community has growing data size problem.
- 3 HPC community has growing need for data analytics.



## Elevating R to Supercomputers

## Elevating R to Supercomputers

- 1 Existing code.

## Elevating R to Supercomputers

- 1 Existing code.
- 2 Syntax.

## Elevating R to Supercomputers

- 1 Existing code.
- 2 Syntax.
- 3 **Philosophy.**

## Programming with Big Data in R (pbdR)

*Productivity, Portability, Performance*





- *Free*<sup>a</sup> R packages.
- Bridging high-performance C with high-productivity of R
- Distributed data details implicitly managed.
- Methods have syntax *identical* to R.

---

<sup>a</sup>MPL, BSD, and GPL licensed

## pbdR Packages



## pbdMPI vs Rmpi

## pbdMPI vs Rmpi

### Reduce Operation with Rmpi

```
1 # int
2 mpi.allreduce(x, type=1)
3 # double
4 mpi.allreduce(x, type=2)
```

### Reduce Operation with pbdMPI

```
1 allreduce(x)
```



## pbdMPI vs Rmpi

### Reduce Operation with Rmpi

```
1 # int
2 mpi.allreduce(x, type=1)
3 # double
4 mpi.allreduce(x, type=2)
```

### Reduce Operation with pbdMPI

```
1 allreduce(x)
```

```
1 > is.integer(1)
2 [1] FALSE
3 > is.integer(2)
4 [1] FALSE
5 > is.integer(1:2)
6 [1] TRUE
```

## pbdMPI vs Rmpi

**Table:** Runtimes (seconds) for  $10,000 \times 10,000$  allgather with **Rmpi** and **pbdMPI**.

Cores	Rmpi	pbdMPI	Speedup
32	24.6	6.7	3.67
64	25.2	7.1	3.55
128	22.3	7.2	3.10
256	22.4	7.1	3.15

## pbdR Example Syntax

```
1 x <- x[-1, 2:5]
2 x <- log(abs(x) + 1)
3 xtx <- t(x) %*% x
4 ans <- svd(solve(xtx))
```


Look familiar?

*The above runs on 1 core with R or 10,000 cores with pbdR*

## Shared and Distributed Memory Machines


### Shared Memory

Direct access to read/change memory (one node)



### Distributed

No direct access to read/change memory.



## Shared and Distributed Memory Machines

### Shared Memory Machines

Thousands of cores



*Nautilus*, University of Tennessee  
1024 cores  
4 TB RAM

### Distributed Memory Machines

Hundreds of thousands of cores



*Kraken*, University of Tennessee  
112,896 cores  
147 TB RAM

# Contents

1 Introduction

2 Benchmarks

3 Challenges

## Non-Optimal Choices Throughout

- 1 Only libre software used (no MKL, ACML, etc.).

## Non-Optimal Choices Throughout

- 1 Only libre software used (no MKL, ACML, etc.).
- 2 1 core = 1 MPI process.



## Non-Optimal Choices Throughout

- 1 Only libre software used (no MKL, ACML, etc.).
- 2 1 core = 1 MPI process.
- 3 No tuning for data distribution.

## Benchmark Data

- 1 Random normal  $N(100, 10000)$ .

## Benchmark Data

- 1 Random normal  $N(100, 10000)$ .
- 2 Local problem size of  $\approx 43.4 \text{ MiB}$ .

## Benchmark Data

- 1 Random normal  $N(100, 10000)$ .
- 2 Local problem size of  $\approx 43.4 \text{ MiB}$ .
- 3 Three sets: 500, 1000, and 2000 columns.


## Benchmark Data

- 1 Random normal  $N(100, 10000)$ .
- 2 Local problem size of  $\approx 43.4 \text{ MiB}$ .
- 3 Three sets: 500, 1000, and 2000 columns.
- 4 Several runs at different core sizes within each set.

## Covariance Code

```
1 x <- ddmatrix("rnorm", nrow=n, ncol=p, mean=mean, sd=sd)
2
3 cov.x <- cov(x)
```

cov()




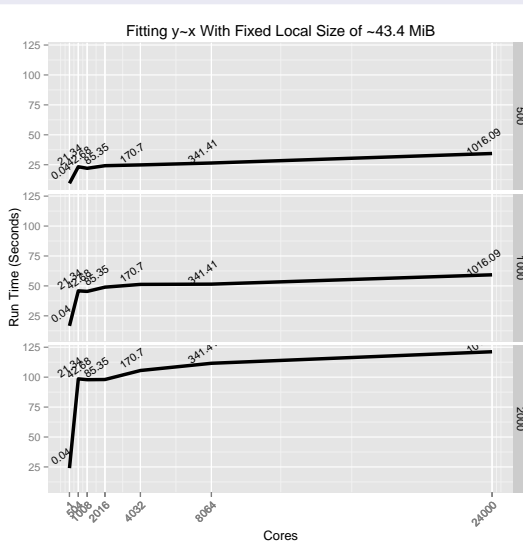
## Linear Model Code


```
1 x <- ddmatrix("rnorm", nrow=n, ncol=p, mean=mean, sd=sd)
2 beta_true <- ddmatrix("runif", nrow=p, ncol=1)
3
4 y <- x %*% beta_true
5
6 beta_est <- lm.fit(x=x, y=y)$coefficients
```



## Data Generation



`lm.fit()`

`lm.fit()`

# Contents

1 Introduction

2 Benchmarks

3 Challenges

## Challenges

- Perceptions.


## Challenges

- Perceptions.
- Library loading.

## Challenges

- Perceptions.
- Library loading.
- Profiling.

## Covariance Revisited: Distributed Data Parameter Calibration





## Tutorials

- XSEDE13, July 22, San Diego, California, USA
- SC13, November 17-22, Denver, Colorado, USA

## Invited Talks

- JSM 2013, August 3-8, Montréal, Québec
- IASC, Aug 22-23, Seoul
- World Statistics Congress, August 25-30, Hong Kong

Thanks for coming!

Questions?

[https://github.com/wrathematics/talks/blob/master/  
user2013/elevatingr.pdf?raw=true](https://github.com/wrathematics/talks/blob/master/user2013/elevatingr.pdf?raw=true)